Asymmetric Errors

Roger Barlow Manchester University, UK and Stanford University, USA

Errors quoted on results are often given in asymmetric form. An account is given of the two ways these can arise in an analysis, and the combination of asymmetric errors is discussed. It is shown that the usual method has no basis and is indeed wrong. For asymmetric systematic errors, a consistent method is given, with detailed examples. For asymmetric statistical errors a general approach is outlined.

1. ASYMMETRIC ERRORS

In the reporting of results from particle physics experiments it is common to see values given with errors with different positive and negative numbers, to denote a 68% central confidence region which is not symmetric about the central estimate. For example (one of many) the Particle Data Group[1] quote

$$B.R.(f_2(1270) \to \pi\pi) = (84.7^{+2.4}_{-1.3})\%.$$

The purpose of this note is to describe how such errors arise and how they can properly be handled, particularly when two contributions are combined. Current practice is to combine such errors separately, i.e. to add the σ^+ values together in quadrature, and then do the same thing for the σ^- values. This is not, to my knowledge, documented anywhere and, as will be shown, is certainly wrong.

There are two separate sources of asymmetry, which unfortunately require different treatments. We call these 'statistical' and 'systematic'; the label is fairly accurate though not entirely so, and they could equally well be called 'frequentist' and 'Bayesian'.

Asymmetric statistical errors arise when the log likelihood curve is not well described by a parabola [2]. The one sigma values (or, equivalently, the 68% central confidence level interval limits) are read off the points at which $\ln L$ falls from its peak by $\frac{1}{2}$ – or, equivalently, when χ^2 rises by 1. This is not strictly accurate, and corrections should be made using Bartlett functions[3], but that lies beyond the scope of this note.

Asymmetric systematic errors arise when the dependence of a result on a 'nuisance parameter' is non-linear. Because the dependence on such parameters – theoretical values, experimental calibration constants, and so forth – is generally complicated, involving Monte Carlo simulation, this study generally has to be performed by evaluating the result x at the $-\sigma$ and $+\sigma$ values of the nuisance parameter a (see [4] for a fuller account) giving σ_x^- and σ_x^+ . $(a \pm \sigma$ gives σ_x^\pm or σ_x^\mp according to the sign of $\frac{dx}{da}$.)

This note summarises a full account of the procedure for asymmetric systematic errors which can be found in [5] and describes what has subsequently been achieved for asymmetric statistical errors. For another critical account see [6].

2. ASYMMETRIC SYSTEMATIC ERRORS

If σ_x^- and σ_x^+ are different then this is a sign that the dependence of x on a is non-linear and the symmetric distribution in a gives an asymmetric distribution in x. In practice, if the difference is not large, one might be well advised to assume a straight line dependence and take the error as symmetric, however we will assume that this is not a case where this is appropriate. We consider cases where a non-linear effect is not small enough to be ignored entirely, but not large enough to justify a long and intensive investigation. Such cases are common enough in practice.

2.1. Models

For simplicity we transform a to the variable u described by a unit Gaussian, and work with X(u) = x(u) - x(0). It is useful to define the mean σ , the difference α , and the asymmetry A:

$$\sigma = \frac{\sigma^+ + \sigma^-}{2} \qquad \alpha = \frac{\sigma^+ - \sigma^-}{2} \qquad A = \frac{\sigma^+ - \sigma^-}{\sigma^+ + \sigma^-}$$
(1)

There are infinitely many non-linear relationships between u and X that will go through the three determined points. We consider two. We make no claim that either of these is 'correct'. But working with asymmetric errors must involve some model of the non-linearity. Practitioners must select one of these two models, or some other (to which the same formalism can be applied), on the basis of their knowledge of the problem, their preference and experience.

Model 1: Two straight lines
 Two straight lines are drawn, meeting at the central value

$$X = \sigma^{+}u \qquad u \ge 0$$

= $\sigma^{-}u \qquad u \le 0.$ (2)

• Model 2: A quadratic function

The parabola through the three points is

$$X = \sigma u + \alpha u^2 = \sigma u + A\sigma u^2. \tag{3}$$

These forms are shown in Figure 1 for a small asymmetry of 0.1, and a larger asymmetry of 0.4.

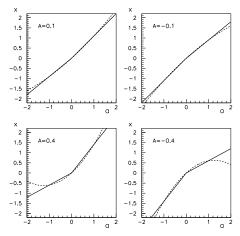


Figure 1: Some nonlinear dependencies

Model 1 is shown as a solid line, and Model 2 is dashed. Both go through the 3 specified points. The differences between them within the range $-1 \le u \le 1$ are not large; outside that range they diverge considerably.

The distribution in u is a unit Gaussian, G(u), and the distribution in X is obtained from $P(X) = \frac{G(u)}{|dX/du|}$. Examples are shown in Figure 2. For Model 1 (again a solid line) this gives a dimidated Gaussian - two Gaussians with different standard deviation for X>0 and X<0. This is sometimes called a 'bifurcated Gaussian', but this is inaccurate. 'Bifurcated' means 'split' in the sense of forked. 'Dimidated' means 'cut in half', with the subsidiary meaning of 'having one part much smaller than the other' [7]. For Model 2 (dashed) with small asymmetries the curve is a distorted Gaussian, given by $\frac{G(u)}{|\sigma+2\alpha u|}$ with $u=\frac{\sqrt{\sigma^2+4\alpha X}-\sigma}{2\alpha}$. For larger asymmetries and/or larger |X| values, the second root also has to be considered.

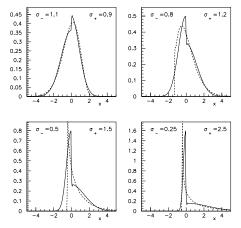


Figure 2: Probability Density Functions from Figure 1

It can be seen that the Model 1 dimidated Gaussian and Model 2 distorted Gaussian are not dissimilar if the asymmetry is small, but are very different if the asymmetry is large.

2.2. Bias

If a nuisance parameter u is distributed with a Gaussian probability distribution, and the quantity X(u) is a nonlinear function of u, then the expectation $\langle X \rangle$ is not $X(\langle u \rangle)$.

For model 1 one has

$$\langle X \rangle = \frac{\sigma^+ - \sigma^-}{\sqrt{2\pi}}$$
 (4)

For model 2 one has

$$\langle X \rangle = \frac{\sigma^{+} - \sigma^{-}}{2} = \alpha \tag{5}$$

Hence in these models, (or any others), if the result quoted is X(0), it is not the mean. It differs from it by an amount of the order of the difference in the positive and negative errors. It is perhaps defensible as a number to quote as the result as it is still the median - there is a 50% chance that the true value is below it and a 50% chance that it is above.

2.3. Adding Errors

If a derived quantity z contains parts from two quantities x and y, so that z = x + y, the distribution in z is given by the convolution:

$$f_z(z) = \int dx f_x(x) f_y(z - x) \tag{6}$$

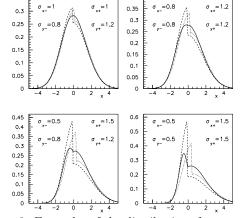


Figure 3: Examples of the distributions from combined asymmetric errors using Model 1.

With Model 1 the convolution can be done analytically. Some results for typical cases are shown in

Figure 3. The solid line shows the convolution, the dashed line is obtained by adding the positive and negative standard deviations separately in quadrature (the 'usual procedure'). The dotted line is described later.

The solid and dashed curves disagree markedly. The 'usual procedure' curve has a larger skew than the convolution. This is obvious. If two distributions with the same asymmetry are added the 'usual procedure' will give a distribution just scaled by $\sqrt{2}$, with the same asymmetry. This violates the Central Limit Theorem, which says that convoluting identical distributions must result in a combined distribution which is more Gaussian, and therefore more symmetric, than its components. This shows that the 'usual procedure' for adding asymmetric errors is inconsistent.

2.4. A consistent addition technique

If a distribution for x is described by some function, $f(x; x_0, \sigma^+, \sigma^-)$, which is a Gaussian transformed according to Model 1 or Model 2 or anything else, then 'combination of errors' involves a convolution of two such functions according to Equation 6. This combined function is not necessarily a function of the same form: it is a special property of the Gaussian that the convolution of two Gaussians gives a third. The (solid line) convolution of two dimidated Gaussians is not itself a dimidated Gaussian. Figure 3 is a demonstration of this.

Although the form of the function is changed by a convolution, some things are preserved. The semi-invariant cumulants of Thièle (the coefficients of the power series expansion of the log of the Fourier Transform) add under convolution. The first two of these are the usual mean and variance. The third is the unnormalised skew:

$$\gamma = \langle x^3 \rangle - 3 \langle x \rangle \langle x^2 \rangle + 2 \langle x \rangle^3$$
 (7)

Within the context of any model, a consistent approach to the combination of errors is to find the mean, variance and skew: μ , V and γ , for each contributing function separately. Adding these up gives the mean, variance and skew of the combined function. Working within the model one then determines the values of σ_-, σ_+ , and x_0 that give this mean, variance and skew.

2.5. Model 1

For Model 1, for which $\langle x^3 \rangle = \frac{2}{\sqrt{2\pi}} (\sigma_+^3 - \sigma_-^3)$ we have

$$\mu = x_0 + \frac{1}{\sqrt{2\pi}} (\sigma^+ - \sigma^-)$$

$$V = \sigma^2 + \alpha^2 \left(1 - \frac{2}{\pi}\right)$$

$$\gamma = \frac{1}{\sqrt{2\pi}} \left[2(\sigma^{+3} - \sigma^{-3}) - \frac{3}{2}(\sigma^+ - \sigma^-)(\sigma^{+2} + \sigma^{-2})\right]$$

$$+\frac{1}{\pi}(\sigma^{+} - \sigma^{-})^{3}$$
 (8)

Given several error contributions the Equations 8 give the cumulants μ , V and γ of each. Adding these up gives the first three cumulants of the combined distribution. Then one can find the set of parameters σ^-, σ^+, x_0 which give these values by using Equations 8 in the other sense.

It is convenient to work with Δ , where Δ is the difference between the final x_0 and the sum of the individual ones. The parameter is needed because of the bias mentioned earlier. Even though each contribution may have $x_0 = 0$, i.e. it describes a spread about the quoted result, it has non-zero μ_i through the bias effect (c.f. Equations 4 and 5). The σ^+ and σ^- of the combined distribution, obtained from the total V and γ , will in general not give the right μ unless a location shift Δ is added. The value of the quoted result will shift.

Recalling section B, for the original distribution one could defend quoting the central value as it was the median, even though it was not the mean. The convoluted distribution not only has a non-zero mean, it also (as can be seen in Figure 3) has non-zero median. If you want to combine asymmetric errors then you have to accept that the quoted value will shift. To make this correction requires a real belief in the asymmetry of the error values. At this point practitioners, unless they are sure that their errors really do have a significant asymmetry, may be persuaded to revert to quoting symmetric errors.

Solving the Equations 8 for σ^-, σ^+ and x_0 given μ , V and γ has to be done numerically. A program for this is available on http://www.slac.stanford.edu/~barlow. Some results are shown in the dotted curve of Figure 3 and Table 1.

Table I Adding errors in Model 1

σ_x^-	σ_x^+	σ_y^-	σ_y^+	σ^{-}	σ^+	Δ
1.0	1.0	0.8	1.2	1.32	1.52	0.08
0.8	1.2	0.8	1.2	1.22	1.61	0.16
0.5	1.5	0.8	1.2	1.09	1.78	0.28
0.5	1.5	0.5	1.5	0.97	1.93	0.41

It is apparent that the dotted curve agrees much better with the solid one than the 'usual procedure' dashed curve does. It is not an exact match, but does an acceptable job given that there are only 3 adjustable parameters in the function. If the shape of the solid curve is to be represented by a dimidated Gaussian, then it is plausible that the dotted curve is the 'best' such representation.

2.6. Model 2

The equivalent of Equations 8 are

$$\mu = x_0 + \alpha$$

$$V = \sigma^2 + 2\alpha^2$$

$$\gamma = 6\sigma^2\alpha + 8\alpha^3$$
(9)

As with Method 1, these are used to find the cumulants of each contributing distribution, which are summed to give the three totals, and then Equation 9 is used again to find the parameters of the distorted Gaussian with this mean, variance and skew. The web program will also do these calculations

Some results are shown in Figure 4 and Table II. The true convolution cannot be done analytically but can be done by a Monte Carlo calculation.

Table II Adding errors in Model 2

σ_x^-	σ_x^+	σ_y^-	σ_y^+	σ^{-}	σ^+	Δ
1.0	1.0	0.8	1.2	1.33	1.54	0.10
0.8	1.2	0.8	1.2	1.25	1.64	0.20
0.5	1.5	0.8	1.2	1.12	1.88	0.35
0.5	1.5	0.5	1.5	1.13	2.07	0.53

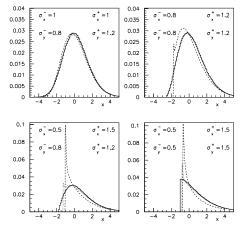


Figure 4: Examples of combined errors using Model 2.

Again the true curves (solid) are not well reproduced by the 'usual procedure' (dashed) but the curves with the correct cumulants (dotted) do a good job. (The sharp behaviour at the edge of the curves is due to the turning point of the parabola.)

2.7. Evaluating χ^2

For Model 1 the χ^2 contribution from a discrepancy δ is just δ^2/σ^{+2} or δ^2/σ^{-2} as appropriate. This is manifestly inelegant, especially for minimisation procedures as the value goes through zero.

For Model 2 one has

$$\delta = \sigma u + A\sigma u^2. \tag{10}$$

This can be considered as a quadratic for u with solution which when squared gives u^2 , the χ^2 contribution, as

$$u^{2} = \frac{2 + 4A\frac{\delta}{\sigma} - 2(1 + 4A\frac{\delta}{\sigma})^{\frac{1}{2}}}{4A^{2}}$$
 (11)

This is not really exact, in that it only takes one branch of the solution, the one approximating to the straight line, and does not consider the extra possibility that the δ value could come from an improbable u value the other side of the turning point of the parabola. Given this imperfection it makes sense to expand the square root as a Taylor series, which, neglecting correction terms above the second power, leads to

$$\chi^2 = \left(\frac{\delta}{\sigma}\right)^2 \left(1 - 2A\left(\frac{\delta}{\sigma}\right) + 5A^2\left(\frac{\delta}{\sigma}\right)^2\right). \tag{12}$$

This provides a sensible form for χ^2 from asymmetric errors. It is important to keep the δ^4 term rather than stopping at δ^3 to ensure χ^2 stays positive! Adding higher orders does not have a great effect. We recommend it for consideration when it is required (e.g. in fitting parton distribution functions) to form a χ^2 from asymmetric errors

2.8. Weighted means

The 'best' estimate (i.e. unbiassed and with smallest variance) from several measurements x_i with different (symmetric) errors σ_i is given by a weighted sum with $w_i = 1/\sigma_i^2$. We wish to find the equivalent for asymmetric errors.

As noted earlier, when sampling from an asymmetric distribution the result is biassed towards the tail. The expectation value $\langle x \rangle$ is not the location parameter x. So for an unbiassed estimator one must take

$$\hat{x} = \sum w_i(x_i - b_i) / \sum w_i \tag{13}$$

where

$$b = \frac{\sigma^{+} - \sigma^{-}}{\sqrt{2\pi}} \quad \text{(Model 1)} \qquad b = \alpha \quad \text{(Model 2)}$$
(14)

The variance of this is given by

$$V = \frac{\sum w_i^2 V_i}{\left(\sum w_i\right)^2} \tag{15}$$

where V_i is the variance of the i^{th} measurement about its mean. Differentiating with respect to w_i to find the minimum gives

$$\frac{2w_i V_i}{\left(\sum w_i\right)^2} - \frac{2\sum w_j^2 V_j}{\left(\sum w_i\right)^3} = 0 \qquad \forall i \tag{16}$$

which is satisfied by $w_i = 1/V_i$. This is the equivalent of the familiar weighting by $1/\sigma^2$. The weights are given, depending on the Model, by (see Equations 8 and 9)

$$V = \sigma^2 + (1 - \frac{2}{\pi})\alpha^2$$
 or $V = \sigma^2 + 2\alpha^2$ (17)

Note that this is not the Maximum Liklelihood estimator - writing down the likelihood in terms of the χ^2 and differentiating does not give a nice form - so in principle there may be better estimators, but they will not have the simple form of a weighted sum.

3. ASYMMETRIC STATISTICAL ERRORS

As explained earlier, (log) likelihood curves are used to obtain the maximum likelihood estimate for a parameter and also the 68% central interval – taken as the values at which $\ln L$ falls by $\frac{1}{2}$ from its peak. For large N this curve is a parabola, but for finite N it is generally asymmetric, and the two points are not equidistant about the peak.

The bias, if any, is not connected to the form of the curve, which is a likelihood and not a pdf. Evaluating a bias is done by integrating over the measured value not the theoretical parameter. We will assume for simplicity that these estimates are bias free. This means that when combining errors there will be no shift of the quoted value.

3.1. Combining asymmetric statistical errors

Suppose estimates \hat{a} and \hat{b} are obtained by this method for variables a and b. a could typically be an estimate of the total number of events in a signal region, and b the (scaled and negated) estimate of background, obtained from a sideband. We are interested in u = a + b, taking $\hat{u} = \hat{a} + \hat{b}$. What are the errors to be quoted on \hat{u} ?

3.2. Likelihood functions known

We first consider the case where the likelihood functions $L_a(\vec{x}|a)$ and $L_b(\vec{x}|b)$ are given.

For the symmetric Gaussian case, the answer is well known. Suppose that the likelihoods are both Gaussian, and further that $\sigma_a = \sigma_b = \sigma$. The log likelihood term

$$\left(\frac{\hat{a}-a}{\sigma}\right)^2 + \left(\frac{\hat{b}-b}{\sigma}\right)^2 \tag{18}$$

can be rewritten

$$\frac{1}{2} \left(\frac{\hat{a} + \hat{b} - (a+b)}{\sigma} \right)^2 + \frac{1}{2} \left(\frac{\hat{a} - \hat{b} - (a-b)}{\sigma} \right)^2 \tag{19}$$

so the likelihood is the product of Gaussians for u = a + b and v = a - b, with standard deviations $\sqrt{2}\sigma$.

Picking a particular value of v, one can then trivially construct the 68% confidence region for u as $[\hat{u}-\sqrt{2}\sigma,\hat{u}+\sqrt{2}\sigma]$. Picking another value of v, indeed any other value of v, one obtains the same region for u. We can therefore say with 68% confidence that these limits enclose the true value of u, whatever the value of v. The uninteresting part of a and b has been 'parametrised away'. This is, of course, the standard result from the combination of errors formula, but derived in a frequentist way using Neyman-style confidence intervals. We could construct the limits on u by finding $\hat{u}+\sigma_u^+$ such that the integrated probability of a result as small as or smaller than the data be 16%, and similarly for σ_u^- , rather than taking the $\Delta \ln L = -\frac{1}{2}$ shortcut, and it would not affect the argument.

The question now is how to generalise this. For this to be possible the likelihood must factorise

$$L(\vec{x}|a,b) = L_u(\vec{x}|u)L_v(\vec{x}|v) \tag{20}$$

with a suitable choice of the parameter v and the functions L_u and L_v . Then we can use the same argument: for any value of v the limits on u are the same, depending only on $L_u(\vec{x}|u)$. Because they are true for any v they are true for all v, and thus in general.

There are cases where this can clearly be done. For two Gaussians with $\sigma_a \neq \sigma_b$ the result is the same as above but with $v = a\sigma_b^2 - b\sigma_a^2$. For two Poisson distributions v is a/b. There are cases (with multiple peaks) where it cannot be done, but let us hope that these are artificially pathological.

On the basis that if it cannot be done, the question is unanswerable, let us assume that it is possible in the case being studied, and see how far we can proceed. Finding the form of v is liable to be difficult, and as it is not actually used in the answer we would like to avoid doing so. The limits on u are read off from the $\Delta \ln L(\vec{x}|u,v) = -\frac{1}{2}$ points where v can have any value provided it is fixed. Let us choose $v = \hat{v}$, the value at the peak. This is the value of v at which $L_v(v)$ is a maximum. Hence when we consider any other value of u, we can find $v = \hat{v}$ by finding the point at which the likelihood is a maximum, varying a - b, or a, or b, or any other combination, always keeping a+bfixed. We can read the limits off a 1 dimensional plot of $\ln L_{max}(\vec{x}|u)$, where the 'max' suffix denotes that at each value of u we search the subspace to pick out the maximum value.

This generalises to more complicated situations. If u = a + b + c we again scan the $\ln L_{max}(\vec{x}|u)$ function, where the subspace is now 2 dimensional.

3.3. Likelihood functions not completely known

In many cases the likelihood functions for a and b will not be given, merely estimates \hat{a} and \hat{b} and their asymmetric errors σ_a^+ , σ_a^- , σ_b^+ and σ_b^- . All we can do is to use these to provide best guess functions $L_a(\vec{x}|a)$ and $L_b(\vec{x}|b)$. A parametrisation of suitable shapes, which for $\sigma^+ \sim \sigma^-$ approximate to a parabola, must be provided. Choosing a suitable parametrisation is not trivial. The obvious choice of introducing small higher-order terms fails as these dominate far from the peak. A likely candidate is:

$$\ln L(a) = -\frac{1}{2} \left(\frac{\ln \left(1 + a/\gamma \right)}{\ln \beta} \right)^2 \tag{21}$$

where $\beta = \sigma_+/\sigma_-$ and $\gamma = \frac{\sigma_+\sigma_-}{\sigma_+-\sigma_-}$. This describes the usual parabola, but with the x-axis stretched by an amount that changes linearly with distance. Figure 5 shows two illustrative results. The first is the Poisson

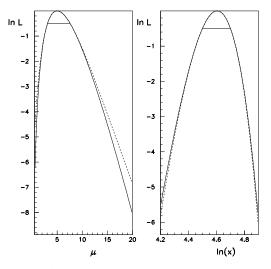


Figure 5: Approximations using Equation 21

likelihood from 5 observed events (solid line) for which the estimate using the $\Delta \ln L = \frac{1}{2}$ points is $\mu = 5^{+2.58}_{-1.92}$, as shown. The dashed line is that obtained inserting these numbers into Equation 21. The second considers a measurement of $x = 100 \pm 10$, of which the logarithm has been taken, to give a value $4.605^{+0.095}_{-0.105}$. Again, the solid line is the true curve and the dashed line the parametrisation. In both cases the agreement is excellent over the range $\approx \pm 1\sigma$ and reasonable over the range $\approx \pm 3\sigma$.

To check the correctness of the method we can use the combination of two Poisson numbers, for which the result is known. First indications are that the errors obtained from the parametrisation are indeed closer to the true Poisson errors than those obtained from the usual technique.

3.4. Combination of Results

A related problem is to find the combined estimate \hat{u} given estimates \hat{a} and \hat{b} (which have asymmetric errors). Here a and b could be results from different channels or different experiments. This can be regarded as a special case, constrained to a=b, i.e. v=0, but this is rather contrived. It is more direct just to say that one uses the log likelihood which is the sum of the two separate functions, and determines the peak and the $\Delta \ln L = -\frac{1}{2}$ points from that. If the functions are known this is unproblematic, if only the errors are given then the same parametrisation technique can be used.

4. CONCLUSIONS

If asymmetric errors cannot be avoided they need careful handling.

A method is suggested and a program provided for combining asymmetric systematic errors. It is not 'rigorously correct' but such perfection is impossible. Unlike the usual method, it is at least open about its assumptions and mathematically consistent.

Formulæ for χ^2 and weighted sums are given.

A method is proposed for combining asymmetric statistical errors if the likelihood functions are known. Work is in progress to enable it to be used given only the results and their errors.

Acknowledgments

The author a gratefully acknowledges the support of the Fulbright Foundation.

References

- [1] D.E. Groom et al., Eur. Phys. J. C15 1 (2000).
- [2] W. T. Eadie et al, "Statistical Methods in Experimental Physics", North Holland, 1971.
- [3] A.G. Frodesen et al. "Probablity and Statistics in Particle Physics", Universitetsforlaget Bergen-Oslo-Tromso (1979), pp 236-239.
- [4] R. J. Barlow "Systematic Errors: Facts and Fictions" in Proc. Durham conference on Advanced Statistical Techniques in Particle Physics, M. R. Whalley and L. Lyons (Eds). IPPP/02/39. 2002.
- [5] R. J. Barlow, "Asymmetric Systematic Errors" preprint MAN/HEP/03/02, ArXiv:physics/030613.
- [6] G. D'Agostini "Bayesian Reasoning in Data Analysis: a Critical Guide", World Scientific (2003).
- [7] The Shorter Oxford English Dictionary, Vol I (A-M) p 190 and p 551 of the 3rd edition (1977).